Limitations of Anonymous Grad Cafe Data

Grad Cafe data is useful for learning SQL because it contains many real-looking application records, but it is not a representative sample of all graduate school applicants. The site depends on anonymous self-submission, so the data can overrepresent people who are especially motivated to report good news, bad news, or unusual outcomes. It can also contain missing values, inconsistent labels, duplicate-looking entries, rounding, exaggeration, or mistakes because there is no admissions office verifying each GPA, GRE score, school name, program, or decision status. Those limits mean the SQL answers should be read as patterns in the submitted Grad Cafe records, not as official statistics about graduate admissions.

The analysis results may look surprising when compared with national or institutional standards because the people who report to Grad Cafe are probably not distributed like the full applicant population. For example, an average GRE quantitative score from Grad Cafe can be much higher than a broader benchmark if applicants with strong scores are more willing to disclose them, if international STEM applicants are overrepresented, or if users without scores leave the field blank and are excluded from averages. In my cleaned dataset, 31,430 of 50,000 records are Fall 2026 entries, about 37.08 percent of those Fall 2026 entries are acceptances, and about 47.73 percent of all rows include comments. The modern-scale GRE quantitative subset averages about 165.81, while the raw GRE field contains mixed scales that make an unfiltered average misleading. Similar selection and formatting effects can affect acceptance rates, GPA averages, and school-specific counts. The data is still valuable for exploratory analysis, but the limitations make careful wording important: the results describe anonymous self-reported entries, not verified admissions outcomes.